


# THE GOOGLE FILE SYSTEM

BY: CHRIS DELLADONNA

# MAIN IDEA:

- The Google file system is a scalable , distributed file system for large distributed data-intensive applications
  - runs on inexpensive commodity hardware
  - provides fault tolerance
  - delivers high aggregate performance to all
- **GFS file system also supports all of their storage needs**
  - Appending data is the focus of performance optimization
  - Block size is 64MB, since they are mostly storing small files
  - Can still effectively store multi-GB files
- **The GFS allows users to work on files concurrently**
  - Users can read, write, and edit at the same time as other users
  - They have prioritized high-sustained bandwidth for processing needs, without emphasis on connection time because of application demand.



# HOW IT IS IMPLEMENTED:


The GFS has been implemented with a lot of specific topics in mind.

- **It is scaled over a lot of inexpensive commodity hardware**
  - Separated into a single-Master and chunk servers over a simplified interface
    - Files are stored Hierarchically
  - Expect various failures and holdups to occur
    - Focus had to be put into error and fault detection
- **GFS provides a lot of space and is extremely scalable**
  - Adding space to chunk servers is very cheap
  - Expected the master could be a potential bottleneck in the system as a result of the amount of chunk pointers that it stores.
    - Each pointer is roughly 64 bits, meaning that the space is accounted for
    - Adding space to the master is also cheap and easy
- **Constant mutations from multiple connections are kept consistent**
  - logging mutations to an operation log , which is stored on the master's local disk and remotely for reliability
- **The system also has the ability to consistently scan through its entire state in the background.**
  - used for chunk garbage collection
  - re-replication in the presence of chunk server failure,
  - chunk migration to balance load and disk space
- **Designed to reduce network overhead**
  - GFS polls chunk servers for that information at startup
  - Master keeps itself up to date by controlling chunk placement and monitoring chunk server status with regular heartbeat messages.

# MY ANALYSIS

**In my opinion, the GFS engineering team has seemed to account for everything that they could experience during the lifespan of the system by:**

- Having fantastic idea how to deal with issues that will more than likely occur resulting from running the system on commodity hardware
- Implementing heartbeat monitors constantly check for failures and interruptions of chunk servers
- Accounting for the amount of space that they should expect to need in the future.
- Realizing the obvious bottlenecks in designing the system the way that they did
  - adding space to the chunk servers is very easy for the degree of service that it provides
- Achieving high reliability by keeping a method to recover information even in catastrophic situations.
  - with the snapshot feature, they always have a checkpoint to look back to
    - if there was an error in creating the checkpoint, that checkpoint will be skipped when the system restarts and it will go back the checkpoint beforehand.






# ADVANTAGES/DISADVANTAGES:

## Advantages

- The file system is very realizable because of the implemented status scanning of the servers involved
- Manages space very efficiently
- Handles traffic based on their needs and their analysis of how most people use their system
- Very fast; the GFS puts the processing load on the host computer and on the chunk servers; very minimal on the masters
- Each chunk server also allows concurrent appends to occur making it very useful for common tasks shared by a group of people.

## Disadvantages:

- Quantity and quality of the components virtually guarantee that some are not functional at any given time
- Upon many people connecting to the same chunk, the chunk can become a hotspot, which can lead to overloading the server with traffic
  - (NOTE: They have fixed this problem for now by storing such executables that affect several chunks at the same time with a higher replication factor and stagger application startup times. However, this is not a long-term solution).



# REAL WORLD CASES:

Real world cases of this technology, I would say, mostly consists of the use from everyday people.

## **Common folk:**

- The average person can effectively access, share, and modify files easily and without high demand on their personal computers.
- They can also store application data for business needs and have the ability to access that information from any host computer that has internet connection.

## **Students:**

- It allows the common student to work on any amounts of projects/schedules/homework/notes on the cloud without having to worry about saving your work or losing what you have done.
- It also makes group work very simple because you can concurrently append documents without getting in the way of other people
  - That file does not have a very limited size and it can also be accessed by any person that has efficient permission to access it.

